**InferPilot**

Agentic inference optimization

# From opaque LLM spend to tested optimization.

For startup AI engineers and CTOs who need to improve cost and latency without quietly sacrificing answer quality.

### THE PRODUCT

Observe real requests → Diagnose

Propose → Human approval

Benchmark → Keep or reject

### WHY NOW

LLM applications are multiplying. Teams need evidence that an optimization helps their own workload.

### FIRST CUSTOMERS

Developer communities → technical design partners → small AI teams. Start with one gateway integration and a reproducible experiment.

---

NovaMaison is the test application. Target: Aiven Runtime + Aiven for PostgreSQL. Gateway metrics + real tools + an isolated exact-cache benchmark. No unmeasured savings claimed.